



Feature Attribution & Interactions

Global & Local Impact




Conclusion:
Orthogonal Feature
Space / No Hidden
Synergies.

Friedman's H-Statistic: Pairwise Feature Interaction Strength




Trustworthy AI: Uncertainty & Fairness

Mathematical Coverage




To maintain a 95% theoretical coverage guarantee, the average prediction set size expands to 1.87.

Epistemic Uncertainty




Monte Carlo Dropout reveals predictive confidence is uniformly uncertain near the decision boundary.

Fairness Audit



Mathematical guarantee of demographic neutrality.


Algorithmic Recourse: Microsoft DiCE



Transitioning from passive classification to algorithmic user agency by mapping minimal counterfactual boundaries.


Causal Uplift Segment Targeting

Causal Uplift Score Distribution (ITE)



Average Treatment Effect: 0.0104. The mathematical reality is that static profile metrics carry no significant causal effect.

App Targeting Segments (Causal Prescriptive)



Engine: T-Learner Meta-Classifler / Bias Elimination: K-Fold Cross-Fitting.

SwipeIQ V2 Dashboard & Recommendations

- 9 Interactive Playgrounds
- Dual-Caching Architecture (<200ms rendering)
- Live Cloud Enterprise Deployment



Pivot Matrix

~~Static Metrics
(Zodiac, Demographics,
App Usage)~~

Dynamic Signals
(NLP Bios, Behavioral
Latency)

A complete **paradigm shift**
away from superficial data.